
Comparing Python KD-Tree Implementations with Focus on Point Cloud Processing

What is the most effective way to process lidar and SfM
point clouds?

Bodo Bookhagen, bodo.bookhagen@uni-potsdam.de,
University of Potsdam

Oct-18-2020

Contents

1	Introduction and Motivation: LidarPC-KDTree	5
2	Environment Installation	5
3	Methods and Approach	6
4	Test datasets	7
4.1	Lidar and SfM Point Clouds from the University of Potsdam Campus Golm	7
4.2	Airborne Lidar data from Santa Cruz Island, California	8
5	Results	8
5.1	Subset of Pozo catchments (n=3,348,668 points)	8
5.1.1	Summary of Results	8
5.1.2	Comparing pyKDTree and cKDTree for 12, 24, and 40 cores	14
5.1.3	Comparing multi-core cKDTree and multi-core FLANN (Fast Library for Approximate Nearest Neighbors) approaches	14
6	Codes	14

List of Figures

1	Generation and query times for single-core <i>sklearnKDTree</i> for varying leafsizes. . . .	9
2	The multi-core <i>ckDTree</i> implementation in <i>scipy.spatial.ckDTree</i> performs well - but you need to set the 'jobs=-1' parameter in the query to achieve best results and use all available cores (only during queries).	10
3	Direct comparison of single core <i>sklearnKDTree</i> and multi core <i>ckDTree</i> approaches for k=5 neighbors. Note the logarithmic y scale - <i>ckDTree</i> is more than one order of magnitude faster.	11
4	Direct comparison of single core <i>sklearnKDTree</i> and multi core <i>ckDTree</i> approaches. .	12
5	Varying leaf size for <i>pyKDTree</i> and <i>ckDTree</i> . The leaf size does not have an significant impact on querying time on multi-core systems (although minor differences can be noted). There is an advantage of multiple cores for higher numbers of neighbors (higher k values). The <i>ckDTree</i> algorithm appears to be the fastest for searches of large k	14

List of Tables

1	List of KD Tree Python implementations	6
2	List of Point clouds from the University of Potsdam Campus Golm.	7
3	Comparison of fastest processing times (any leaf size) for all implemented algorithms in seconds. Note that the scipy standard <i>KDTree</i> has not been processed due to the very slow processing times. All times are the average of 3 iterations.	12
4	Worst (slowest times in seconds) for any processing steps - mostly reflecting the impact of leaf sizes.	13
5	Range of times between different leaf sizes. Most notable differences are for the single core <i>sklearnKDTree</i> approach.	13
6	Best leaf sizes (fastest times). Note the differences for varying numbers of neighbors. .	13

1 Introduction and Motivation: LidarPC-KDTree

Comparison of KDTree implementations for Lidar PointClouds (PC) and Structure-from-Motion (SfM) dataset.

One of the core processing steps for irregular PC is to understand the neighborhood for each point (kNN - k-Nearest-Neighbors). This is often done using [KD Trees](#). There exist myriad of implementations for various applications and KD Trees have become an important tool for Deep Learning that have been implemented in [kNN \(k-nearest neighbor\)](#) algorithms. Many of the approaches have been optimized for multi-dimensional datasets ($n > 5$ and up to 50). In the recent months and years, KD-Trees relying on [CUDA](#) or [OpenCL](#) implementations have become more coming and easily approachable through Python or Matlab interfaces.

Here, we briefly explore existing algorithms and test, which one perform favorable for 3D lidar PC (or SfM PC). We only use three dimensions (x, y, z), but future implementation may rely on four (intensity) or higher dimensional lidar PC. We focus on implementations accessible through Python or C.

We note that there exist other algorithm and parameter comparisons (e.g. [knn-benchmarking in python](#) and [knn-benchmarking](#)) and these are very useful and helpful – but these are neither tailored for lidar/SfM PC nor have been using recent implementations. Most comparison also focus on the general applicability of KD-Tree algorithm and explore the impact of leaf sizes and dimensionality - both parameters do not change for lidar PC.

2 Environment Installation

See [miniconda installation instructions](#) to setup an environment for processing. The installation of the python codes is done through conda on a Ubuntu 18.04 LTS system (also tested on 20.04 LTS).

Conda installation:

```
1 conda create -y -n PC_py3 -c anaconda -c conda-forge -c defaults ipython spyder \
   python=3.8 gdal=3 numpy scipy dask h5py pandas pytables hdf5 cython matplotlib \
   tabulate scikit-learn pykdtree pyflann cyflann scikit-image opencv ipywidgets \
   scikit-learn gmt=6* imagemagick
```

Next:

```
1 conda activate PC_py3
2 pip install laspy
3 pip install tables
```

3 Methods and Approach

We construct the following scenarios: 1. Deriving $k=5,10,50,100,500,1000$ nearest neighbors from four lidar/SfM point clouds with 14e6, 38e6, 69e6, and 232e6 (million) points. 2. We time the generation of a KD-Tree and the queries separately for each. 3. Searching for neighbors within a given search radius/sphere (not supported by all algorithms). 4. The k -nearest neighbors can be used to estimate point-density or perform further classification on the neighborhood structure of points (e.g., curvature)

We note that we query the tree with all points (e.g., $k=50$ neighbors for all points) and thus create large queries for neighborhood statistical analysis.

An incomplete list of available algorithms and implementations. *Note: We have not used all of them for the tests, because some implementations are very slow and mostly for instructive/teaching purposes.* Also, in all instances we have used the standard options and parameters, but these may not always be the most useful ones.

Table 1: List of KD Tree Python implementations

Name	Reference and Documentation	Comments
scipy.spatial.KDTree	Manual	Pure Python implementation of KD tree. Querying is very slow and usage is not suggested. Not used. <i>Single core CPU processing.</i>
scipy.spatial.cKDTree	Manual	KDTree implementation in Cython. <i>Single and Multi-core CPU processing.</i>
sklearn.neighbors.KDTree	Manual	KDTree implementation in sklearn. <i>Single core CPU processing.</i>
pyKDTree	github page and pypi project page	fast implementation for common use cases (low dimensions and low number of neighbours) for both tree construction and queries. The implementation is based on scipy.spatial.cKDTree and libANN by combining the best features from both and focus on implementation efficiency. <i>Multi-core CPU processing.</i>
pyflann	github	pyflann is the python bindings for FLANN - Fast Library for Approximate Nearest Neighbors and FLANN Manual 1.8.4 <i>Multi-core CPU processing.</i>

Name	Reference and Documentation	Comments
cyflann	github	cyflann is the a cython interface for FLANN - Fast Library for Approximate Nearest Neighbors and FLANN Manual 1.8.4 . <i>Multi-core CPU processing.</i>
NearestNeighbors	cuml-kNN	GPU implementaiton of knn via rapidsai , <i>currently only supports brute-force algorithm and is not competitive</i>

4 Test datasets

4.1 Lidar and SfM Point Clouds from the University of Potsdam Campus Golm

The dense and high-resolution point clouds have been generated between 2018-2020 for parts of the University of Potsdam Campus Golm. These represent mixed-urban environments with building and vegetation. The files are too large for github and have been stored on Dropbox (links provided below).

Table 2: List of Point clouds from the University of Potsdam Campus Golm.

Name	PC Type	# of points	Point Density [pts/m2]
Golm_May06_2018	Airborne Lidar	14,437,532	61
mavicpro2	Mavic Pro2, Agisoft Photoscan, high quality processing setting, images from nadiar and 15 degree angle taken	38,334,551	219
mavicpro2_06Sept2019	Mavic Pro2, Agisoft Photoscan, high quality processing setting	69,482,218	707
inspire2	Inspire 2, high quality processing, 1031 images	232,269,911	988

4.2 Airborne Lidar data from Santa Cruz Island, California

The airborne point cloud from Santa Cruz Island, California represents a natural terrain without buildings, but lower density. The dataset contains 3,348,668 points with a point-density of 7.2 pts/m² and has been ground-classified using [LAStools](#). The points have been colored using an airphoto from the same time as the lidar flight. The test area is from a small subset of the Pozo catchment in the southwestern part of the island. These data are openly accessible and available from [opentopography](#) and were originally acquired by the USGS in 2010. The geologic and geomorphic environment and setting of the Santa Cruz Island has been described in several peer-reviewed scientific publications (e.g., [Perroy et al., 2010](#), [Perroy et al., 2012](#), [Neely et al., 2017](#), [Rheinwalt et al., 2019](#), [Clubb et al., 2019](#)).

5 Results

We ran tests on a AMD Ryzen Threadripper 2970WX 24-Core Processor (2019) with a NVIDIA Titan RTX 24 GB running Ubuntu 18.04 (CUDA 10.1) and a AMD Ryzen 9 3900X 12-Core Processor with a NVIDIA GeForce RTX 2080 SUPER running Ubuntu 20.04 (CUDA 11.0).

5.1 Subset of Pozo catchments (n=3,348,668 points)

A first test using standard single-core and multi-core algorithms for n=3,348,668 queries for n=3,348,668 points. Note that the KDTree calculations from *scipy.spatial.KDTree* have not been included, because they were too slow. Also, for the single-core *sklearnKDTree* approach, no higher number of neighbors have been included (too slow). All results show times in seconds (s) and have been averaged over n=3 runs.

5.1.1 Summary of Results

Comparing the traditional and widely used *sklearnKDTree* (single core), *cKDTree* (multi core), *pyKDTree* (multi core), and FLANN approaches we note the following results: 1. The leaf size is an important parameter to speed up single-core querying trees. Depending on point cloud structure, different leaf sizes provide very different results and can improve query times. We note that the default leaf size does not generate useful results for real-world airborne lidar data and that there exists a minimum time representing an optimal leaf size (cf. Figure 1). 2. The *sklearnKDTree* (single core) is slow on these massive queries. The option `dualtree=True` has been used to speed up processing. 3. *cKDTree* with `jobs=-1` set for querying outperforms single-core approaches - especially on modern multi-core systems. Leaf size does not have a significant impact on multi-core processing, but some for larger neighborhood queries (k>100) (cf. Figures 2 and {pc_sklearnKDTree_cKDTree_k5_AMD3900X_12cores}). 4. There are minimal difference between different approach. For example. the max. difference between

sklearnKDTree and *cKDTree* is 0.2m and the median difference is 0.0 (see Figure 4). 5. Comparing *cKDTree* with 12, 24, and 48 core processors indicates a clear advantage of multi-threading processes. (cf. Figure 5). We emphasize that in order to take full advantage of multi-threading processes, an increase in available DRAM is needed (i.e., more cores require more DRAM). We note that *pyKDTree* has lower peak memory requirement than *cKDTree*. 6. The FLANN (Fast Library for Approximate Nearest Neighbors) family of approaches provides additional advancements, especially for large datasets and massive queries. 7. Initial tests with cuML (CUDA RAPIDS) show that the implemented brute-force approach for nearest neighbor searches is not competitive against the multi-core approaches (*cKDTree* and *pyKDTree*) and highly optimized FLANN approaches. But there are other processing advantages of data analysis using CUDA Dataframes (cudf).

sklearnKDTree (single core): generation and query (AMD3900X: 12 cores)

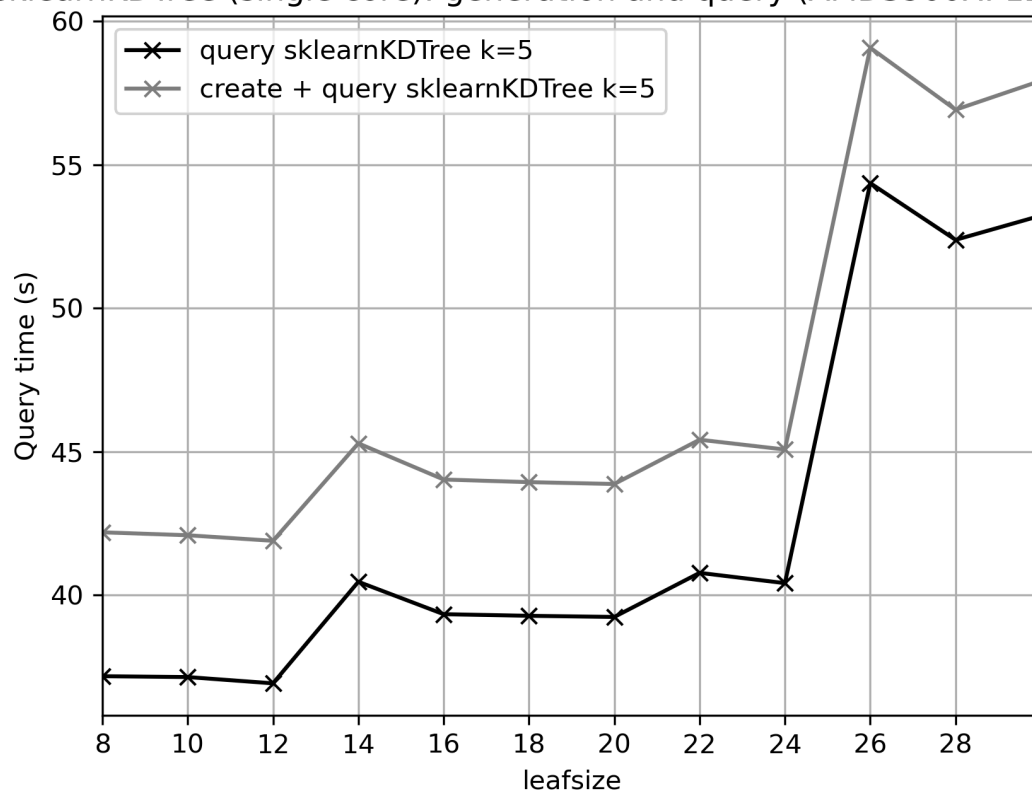


Figure 1: Generation and query times for single-core sklearnKDTree for varying leafsizes.

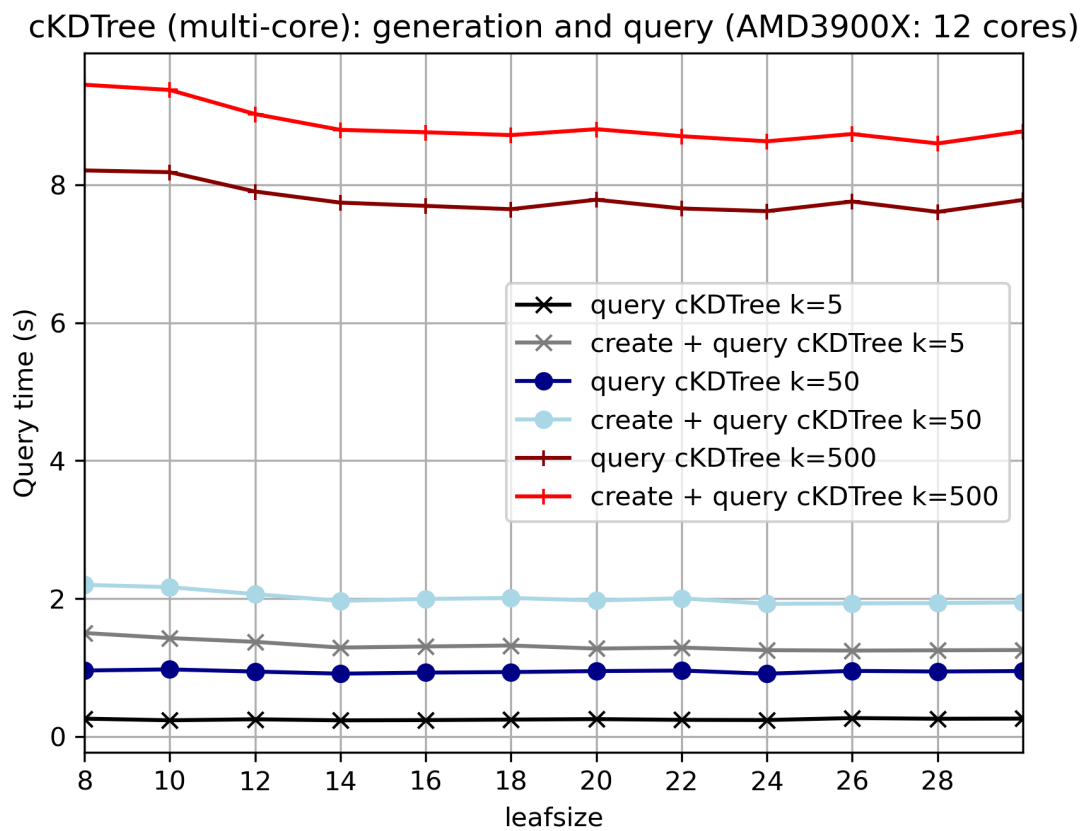


Figure 2: The multi-core cKDTree implementation in *scipy.spatial.cKDTree* performs well - but you need to set the 'jobs=-1' parameter in the query to achieve best results and use all available cores (only during queries).

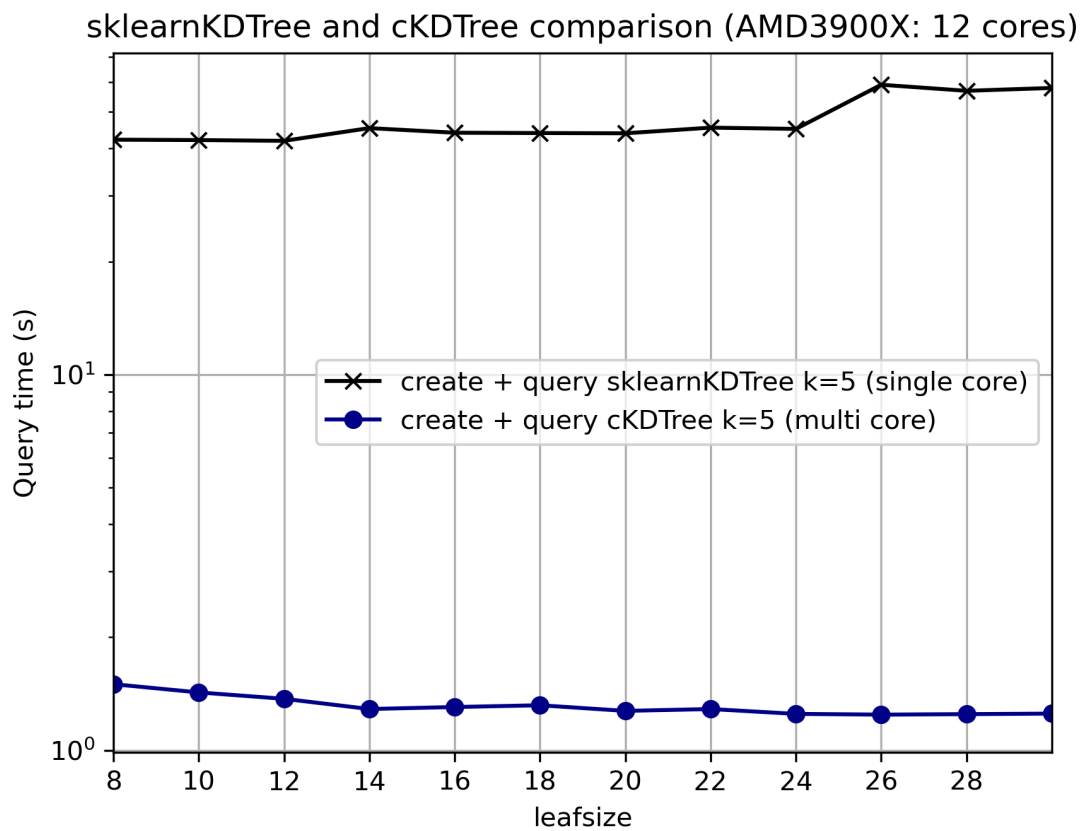


Figure 3: Direct comparison of single core *sklearnKDTree* and multi core *cKDTree* approaches for $k=5$ neighbors. Note the logarithmic y scale - *cKDTree* is more than one order of magnitude faster.

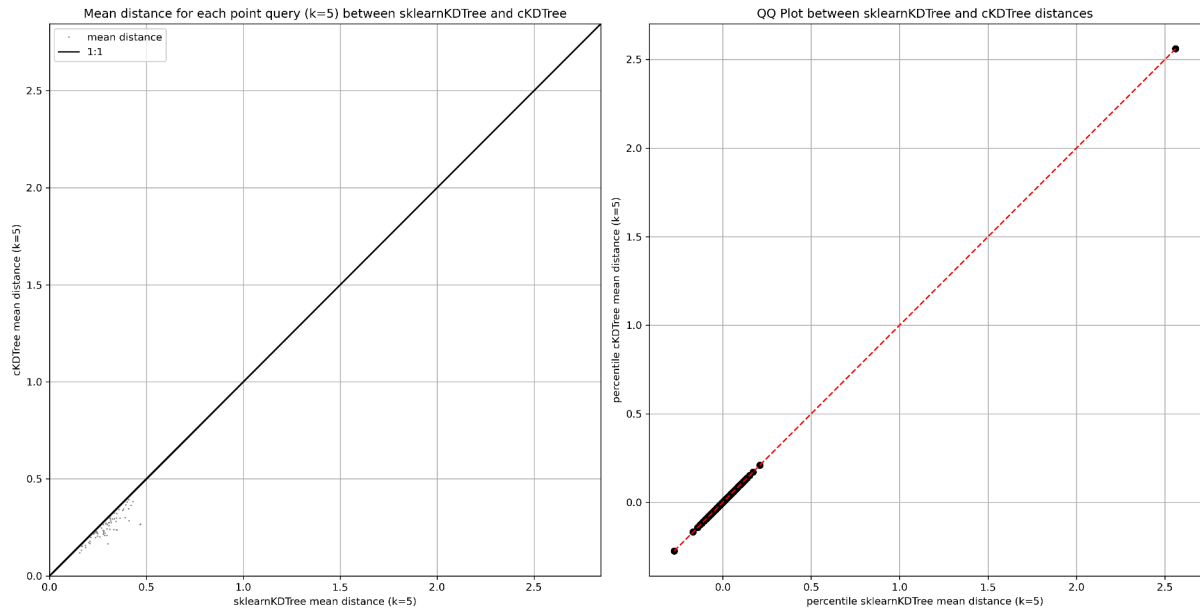


Figure 4: Direct comparison of single core *sklearnKDTree* and multi core *cKDTree* approaches.

Table 3: Comparison of fastest processing times (any leaf size) for all implemented algorithms in seconds. Note that the *scipy* standard *KDTree* has not been processed due to the very slow processing times. All times are the average of 3 iterations.

Algorithm	Generate KDTree (s)	Query k=5 (s)	Query k=10 (s)	Query k=50 (s)	Query k=100 (s)	Query k=500 (s)	Query k=1000 (s)
KDTree	5.25	nan	nan	nan	nan	nan	nan
sklearnKDTree	1.51	36.93	nan	nan	nan	nan	nan
cKDTree	0.32	0.23	0.31	0.91	1.67	7.47	15.13
pyKDTree	0.07	0.23	0.32	1.1	2.58	35.17	129.81
pyflannKDTree	0.19	0.17	0.24	0.97	2.11	12.54	27.4
cyflannKDTree	0.26	0.2	0.26	1	2.2	9.69	20.01

Table 4: Worst (slowest times in seconds) for any processing steps - mostly reflecting the impact of leaf sizes.

Algorithm	Generate KDTree (s)	Query k=5 (s)	Query k=10 (s)	Query k=50 (s)	Query k=100 (s)	Query k=500 (s)	Query k=1000 (s)
KDTree	5.25	nan	nan	nan	nan	nan	nan
sklearnKDTree	1.67	54.35	nan	nan	nan	nan	nan
cKDTree	0.41	0.27	0.39	0.97	1.78	8.21	17.64
pyKDTree	0.12	0.26	0.34	1.17	2.69	35.78	131.58
pyflannKDTree	0.19	0.17	0.24	0.97	2.11	12.54	27.4
cyflannKDTree	0.26	0.2	0.26	1	2.2	9.69	20.01

Table 5: Range of times between different leaf sizes. Most notable differences are for the single core *sklearnKDTree* approach.

Algorithm	Generate KDTree (s)	Query k=5 (s)	Query k=10 (s)	Query k=50 (s)	Query k=100 (s)	Query k=500 (s)	Query k=1000 (s)
sklearnKDTree	0.16	17.43	nan	nan	nan	nan	nan
cKDTree	0.1	0.03	0.08	0.06	0.12	0.74	2.51
pyKDTree	0.05	0.02	0.01	0.07	0.12	0.61	1.77

Table 6: Best leaf sizes (fastest times). Note the differences for varying numbers of neighbors.

Algorithm	Generate KDTree (leaf size)	Query k=5 (leaf size)	Query k=10 (leaf size)	Query k=50 (leaf size)	Query k=100 (leaf size)	Query k=500 (leaf size)	Query k=1000 (leaf size)
KDTree	10	8	8	8	8	8	8
cKDTree	36	14	16	24	14	38	38
sklearnKDTree	28	12	8	8	8	8	8
pyKDTree	36	16	20	16	28	26	32

5.1.2 Comparing pyKDTree and cKDTree for 12, 24, and 40 cores

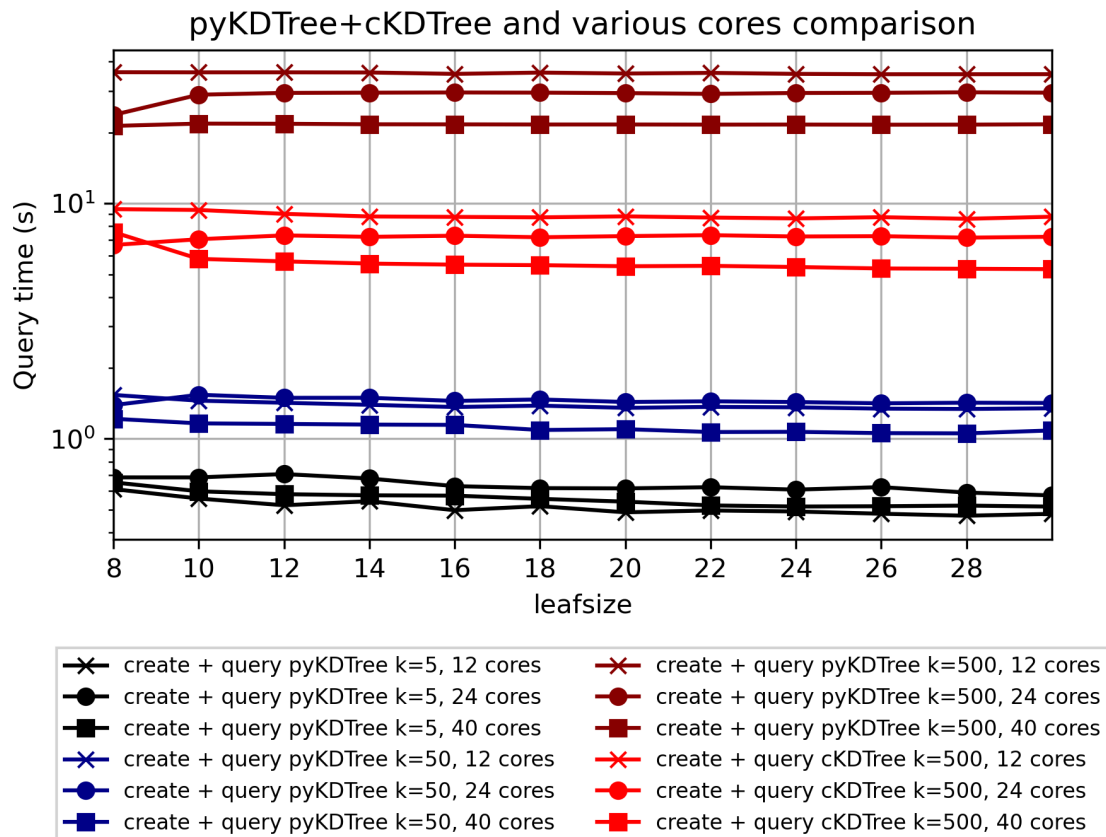


Figure 5: Varying leaf size for *pyKDTree* and *cKDTree*. The leaf size does not have an significant impact on querying time on multi-core systems (although minor differences can be noted). There is an advantage of multiple cores for higher numbers of neighbors (higher *k* values). The *cKDTree* algorithm appears to be the fastest for searches of large *k*

5.1.3 Comparing multi-core cKDTree and multi-core FLANN (Fast Library for Approximate Nearest Neighbors) approaches

6 Codes

All Python codes are available on the github repository [LidarPC-KDTree](#).

We setup separate functions for the generation of KDTrees:

```
1 def pc_generate_KDTree(pc_xyz, leafsize=10):
2     try:
3         from scipy import spatial
```

```
4     except ImportError:
5         raise pc_generate_KDTree("scipy not installed.")
6     pc_xyz_KDTree_tree = spatial.KDTree(pc_xyz, leafsize=leafsizei)
7     return pc_xyz_KDTree_tree
8
9 def pc_query_KDTree(pc_xyz_KDTree_tree, pc_xyz, k=10):
10     pc_kdtree_distance, pc_kdtree_id = pc_xyz_KDTree_tree.query(pc_xyz, k=k)
11     return pc_kdtree_distance, pc_kdtree_id
12
13 def pc_generate_sklearnKDTree(pc_xyz, leafsizei=10):
14     #conda install scikit-learn
15     try:
16         from sklearn.neighbors import KDTree as sklearnKDTree
17     except ImportError:
18         raise pc_generate_sklearnKDTree("sklearn not installed.")
19     pc_xyz_sklearnKDTree_tree = sklearnKDTree(pc_xyz, leaf_size=leafsizei)
20     return pc_xyz_sklearnKDTree_tree
21
22 def pc_query_sklearnKDTree(pc_xyz_sklearnKDTree_tree, pc_xyz, k=10):
23     pc_sklearnKDTree_distance, pc_sklearnKDTree_id = \
24         pc_xyz_sklearnKDTree_tree.query(pc_xyz, k=k, dualtree=True)
25     return pc_sklearnKDTree_distance, pc_sklearnKDTree_id
26
27 def pc_generate_ckDTree(pc_xyz, leafsizei=10):
28     try:
29         from scipy import spatial
30     except ImportError:
31         raise pc_generate_ckDTree("scipy not installed.")
32     pc_xyz_ckDTree_tree = spatial.cKDTree(pc_xyz, leafsize=leafsizei)
33     return pc_xyz_ckDTree_tree
34
35 def pc_query_ckDTree(pc_xyz_ckDTree_tree, pc_xyz, k=10):
36     pc_ckDTree_distance, pc_ckDTree_id = pc_xyz_ckDTree_tree.query(pc_xyz, k=k, \
37         n_jobs=-1)
38     return pc_ckDTree_distance, pc_ckDTree_id
39
40 def pc_generate_pyKDTree(pc_xyz, leafsizei=10):
41     try:
42         from pykdtree.kdtree import KDTree as pyKDTree
43     except ImportError:
44         raise pc_generate_pyKDTree("pykdtree not installed.")
45     pc_xyz_pyKDTree_tree = pyKDTree(pc_xyz, leafsize=leafsizei)
46     return pc_xyz_pyKDTree_tree
47
48 def pc_query_pyKDTree(pc_xyz_pyKDTree_tree, pc_xyz, k=10):
49     pc_pyKDTree_distance, pc_pyKDTree_id = pc_xyz_pyKDTree_tree.query(pc_xyz, k=k)
50     return pc_pyKDTree_distance, pc_pyKDTree_id
51
52 def pc_generate_pyflannKDTree(pc_xyz):
53     #conda install -y -c conda-forge pyflann
54     try:
55         import pyflann
56     except ImportError:
```

```

55     raise pc_generate_pyflannKdTree("pyflann not installed.")
56     pyflann.set_distance_type('euclidean')
57     pc_xyz_pyflannKdTree_tree = pyflann.FLANN()
58     pc_xyz_pyflannKdTree_tree.build_index(pc_xyz, algorithm='kdtree_single', trees=8)
59     return pc_xyz_pyflannKdTree_tree
60
61 def pc_query_pyflannKdTree(pc_xyz_pyflannKdTree_tree, pc_xyz, k=10):
62     pc_pyflannKdTree_id, pc_pyflannKdTree_distance = \
63         pc_xyz_pyflannKdTree_tree.nn_index(pc_xyz, k)
64     return pc_pyflannKdTree_distance, pc_pyflannKdTree_id
65
66 def pc_generate_cyflannKdTree(pc_xyz):
67     #conda install -y -c conda-forge cyflann
68     try:
69         import cyflann
70     except ImportError:
71         raise pc_generate_cyflannKdTree("pyflann not installed.")
72     cyflann.set_distance_type('euclidean')
73     pc_xyz_cyflannKdTree_tree = cyflann.FLANNIndex()
74     pc_xyz_cyflannKdTree_tree.build_index(pc_xyz, algorithm='kdtree_single', trees=8)
75     return pc_xyz_cyflannKdTree_tree
76
77 def pc_query_cyflannKdTree(pc_xyz_cyflannKdTree_tree, pc_xyz, k=10):
78     pc_cyflannKdTree_id, pc_cyflannKdTree_distance = \
79         pc_xyz_cyflannKdTree_tree.nn_index(pc_xyz, k)
80     return pc_cyflannKdTree_distance, pc_cyflannKdTree_id

```

Next, we setup wrapping functions for the timer:

```

1 def wrapper(func, *args, **kwargs):
2     def wrapped():
3         return func(*args, **kwargs)
4     return wrapped

```

And function for loading in LAS/LAZ files (if laszip is installed):

```

1 def load_LAS(las_fname, dtype='float32'):
2     """
3     Load LAS or LAZ file (only coordinates) and return pc_xyz and xy vectors. \
4     Converts float64 to float32 by default, unless you set dtype='float64'
5     """
6     from laspy.file import File
7     inFile = File(las_fname, mode='r')
8     pc_pc_xyz = \
9         np.vstack((inFile.get_x()*inFile.header.scale[0]+inFile.header.offset[0], \
10                    inFile.get_y()*inFile.header.scale[1]+inFile.header.offset[1], \
11                    inFile.get_z()*inFile.header.scale[2]+inFile.header.offset[2])).transpose()
12
13     #setting datatype to float32 to save memory.
14     if dtype == 'float32':
15         pc_pc_xyz = pc_pc_xyz.astype('float32')
16     return pc_pc_xyz

```


Additionally, if you want to run this from the command line:

```

1 def cmdLineParser():
2     parser = argparse.ArgumentParser(description='Compare KDTree algorithms for \
        lidar or SfM PointClouds (PC). B. Bookhagen \
        (bodo.bookhagen@uni-potsdam.de), Aug 2019.')
3     parser.add_argument('-i', '--inlas', type=str, \
        default='Pozo_WestCanada_clg.laz', help='LAS/LAZ file with point-cloud \
        data.')
4     parser.add_argument('--nr_of_repetitions', type=int, default=5, help='Number \
        of repetitions')
5     parser.add_argument('--hdf_filename', type=str, \
        default='Pozo_WestCanada_clg_kdresults_5rep.hdf', help='Output HDF file \
        containing results from iterations.')
6     parser.add_argument('--csv_filename', type=str, \
        default='Pozo_WestCanada_clg_kdresults_5rep.csv', help='Output CSV file \
        containing results from iterations.')
7     parser.add_argument('--nr_of_repetitions_generate', type=int, default=1, \
        help='Number of repetitions used for generating index. Set to 1 to avoid \
        caching effects.')
8     return parser.parse_args()

```

We load the LAS pointcloud:

```

1 print('Loading input file: %s... '%inps.inlas, end='', flush=True)
2 pc_xyz = load_LAS(inps.inlas, dtype='float32')
3 print('loaded %s points'%inps.inlas, format(pc_xyz.shape[0]))

```

And run some iterations for *sklearn*:

```

1 #Run sklearnKDTree
2 pc_generate_sklearnKDTree_time = np.empty( (len(leafrange), 1) )
3 pc_generate_sklearnKDTree_time[:] = np.nan
4 pc_query_sklearnKDTree_k5_time = np.empty( (len(leafrange), 1) )
5 pc_query_sklearnKDTree_k5_time[:] = np.nan
6 pc_query_sklearnKDTree_k10_time = np.empty( (len(leafrange), 1) )
7 pc_query_sklearnKDTree_k10_time[:] = np.nan
8 pc_query_sklearnKDTree_k50_time = np.empty( (len(leafrange), 1) )
9 pc_query_sklearnKDTree_k50_time[:] = np.nan
10 pc_query_sklearnKDTree_k100_time = np.empty( (len(leafrange), 1) )
11 pc_query_sklearnKDTree_k100_time[:] = np.nan
12 pc_query_sklearnKDTree_k500_time = np.empty( (len(leafrange), 1) )
13 pc_query_sklearnKDTree_k500_time[:] = np.nan
14 pc_query_sklearnKDTree_k1000_time = np.empty( (len(leafrange), 1) )
15 pc_query_sklearnKDTree_k1000_time[:] = np.nan
16 pc_query_sklearnKDTree_k5_stats = np.empty( (len(leafrange), len(pc_xyz), 5) )
17 pc_query_sklearnKDTree_k5_stats[:] = np.nan
18 pc_query_sklearnKDTree_k10_stats = np.empty( (len(leafrange), len(pc_xyz), 5) )
19 pc_query_sklearnKDTree_k10_stats[:] = np.nan
20
21 leaf_counter = 0

```

```

22 for leafsizei in leafrange:
23     print('\n\tGenerating sklearnKDTree... with leafsize = %d (%dx)'%(leafsizei, \
        inps.nr_of_repetitions_generate), end='', flush=True)
24     wrapped = wrapper(pc_generate_sklearnKDTree, pc_xyz, leafsizei)
25     pc_generate_sklearnKDTree_time[leaf_counter] = timeit.timeit(wrapped, \
        number=inps.nr_of_repetitions_generate)
26     pc_xyz_sklearnKDTree_tree = pc_generate_sklearnKDTree(pc_xyz, \
        leafsizei=leafsizei)
27     print('time (average of %d runs): %0.3fs or \
        %0.2fm'%(inps.nr_of_repetitions_generate, \
        pc_generate_sklearnKDTree_time[leaf_counter]/inps.nr_of_repetitions, \
        pc_generate_sklearnKDTree_time[leaf_counter]/inps.nr_of_repetitions/60))
28
29     print('\tQuerying sklearnKDTree with k=5 for all points (%dx)... \
       '%(inps.nr_of_repetitions), end='', flush=True)
30     wrapped = wrapper(pc_query_sklearnKDTree, pc_xyz_sklearnKDTree_tree, pc_xyz, k=5)
31     pc_query_sklearnKDTree_k5_time[leaf_counter] = timeit.timeit(wrapped, \
        number=inps.nr_of_repetitions)
32     print('time (average of %d runs): %0.3fs or %0.2fm'%(inps.nr_of_repetitions, \
        pc_query_sklearnKDTree_k5_time[leaf_counter]/inps.nr_of_repetitions, \
        pc_query_sklearnKDTree_k5_time[leaf_counter]/inps.nr_of_repetitions/60))
33     pc_sklearnKDTree_distance, pc_sklearnKDTree_id = \
        pc_query_sklearnKDTree(pc_xyz_sklearnKDTree_tree, pc_xyz, k=5)
34     pc_query_sklearnKDTree_k5_stats[leaf_counter, :, 0] = \
        np.mean(pc_sklearnKDTree_distance, axis=1)
35     pc_query_sklearnKDTree_k5_stats[leaf_counter, :, 1] = \
        np.std(pc_sklearnKDTree_distance, axis=1)
36     pc_query_sklearnKDTree_k5_stats[leaf_counter, :, 2] = \
        np.percentile(pc_sklearnKDTree_distance, [25], axis=1)
37     pc_query_sklearnKDTree_k5_stats[leaf_counter, :, 3] = \
        np.percentile(pc_sklearnKDTree_distance, [50], axis=1)
38     pc_query_sklearnKDTree_k5_stats[leaf_counter, :, 4] = \
        np.percentile(pc_sklearnKDTree_distance, [75], axis=1)
39     pc_sklearnKDTree_distance = None
40     pc_sklearnKDTree_id = None
41
42     pc_xyz_sklearnKDTree_tree = None
43     leaf_counter += 1

```

Results can be plotted with:

```

1 plt.clf()
2 plt.plot(leafrange, pc_query_sklearnKDTree_k5_time/inps.nr_of_repetitions, 'kx-', \
    label='query sklearnKDTree k=5')
3 plt.plot(leafrange, \
    (pc_query_sklearnKDTree_k5_time)/inps.nr_of_repetitions+pc_generate_sklearnKDTree_time/inps.nr_o
    'x-', c='gray', label='create + query sklearnKDTree k=5')
4 plt.title('sklearnKDTree (single core): generation and query (AMD3900X: 12 cores)')
5 plt.grid()
6 plt.xlabel('leafsize')
7 plt.ylabel('Query time (s)')
8 plt.xlim([8,30])
9 plt.xticks(np.arange(8,30,step=2))

```

```

10 plt.ylim([0,2])
11 plt.legend()
12 plt.savefig('figs/pc_sklearnKDTree_AMD3900X_12cores.png', dpi=300, \
    orientation='landscape')

```

For *ckDTree* we perform a more exhaustive analysis:

```

1 #Run ckDTree (cython implementation from scipy)
2 pc_generate_ckDTree_time = np.empty( (len(leafrange), 1) )
3 pc_generate_ckDTree_time[:] = np.nan
4 pc_query_ckDTree_k5_time = np.empty( (len(leafrange), 1) )
5 pc_query_ckDTree_k5_time[:] = np.nan
6 pc_query_ckDTree_k10_time = np.empty( (len(leafrange), 1) )
7 pc_query_ckDTree_k10_time[:] = np.nan
8 pc_query_ckDTree_k50_time = np.empty( (len(leafrange), 1) )
9 pc_query_ckDTree_k50_time[:] = np.nan
10 pc_query_ckDTree_k100_time = np.empty( (len(leafrange), 1) )
11 pc_query_ckDTree_k100_time[:] = np.nan
12 pc_query_ckDTree_k500_time = np.empty( (len(leafrange), 1) )
13 pc_query_ckDTree_k500_time[:] = np.nan
14 pc_query_ckDTree_k1000_time = np.empty( (len(leafrange), 1) )
15 pc_query_ckDTree_k1000_time[:] = np.nan
16 pc_query_ckDTree_k5_stats = np.empty( (len(leafrange), len(pc_xyz), 5) )
17 pc_query_ckDTree_k5_stats[:] = np.nan
18 pc_query_ckDTree_k10_stats = np.empty( (len(leafrange), len(pc_xyz), 5) )
19 pc_query_ckDTree_k10_stats[:] = np.nan
20
21 leaf_counter = 0
22 for leafsizei in leafrange:
23     print('\n\tGenerating ckDTree with leafsize = %d ... (%dx)'%(leafsizei, \
        inps.nr_of_repetitions_generate), end='', flush=True)
24     wrapped = wrapper(pc_generate_ckDTree, pc_xyz, leafsizei)
25     pc_generate_ckDTree_time[leaf_counter] = timeit.timeit(wrapped, \
        number=inps.nr_of_repetitions_generate)
26     pc_xyz_ckDTree_tree = pc_generate_ckDTree(pc_xyz, leafsizei=leafsizei)
27     print('time (average of %d runs): %0.3fs or \
        %0.2fm'%(inps.nr_of_repetitions_generate, \
        pc_generate_ckDTree_time[leaf_counter]/inps.nr_of_repetitions, \
        pc_generate_ckDTree_time[leaf_counter]/inps.nr_of_repetitions/60))
28
29     print('\tQuerying ckDTree with k=5 for all points (%dx)... \
        '%(inps.nr_of_repetitions), end='', flush=True)
30     wrapped = wrapper(pc_query_ckDTree, pc_xyz_ckDTree_tree, pc_xyz, k=5)
31     pc_query_ckDTree_k5_time[leaf_counter] = timeit.timeit(wrapped, \
        number=inps.nr_of_repetitions)
32     print('time (average of %d runs): %0.3fs or %0.2fm'%(inps.nr_of_repetitions, \
        pc_query_ckDTree_k5_time[leaf_counter]/inps.nr_of_repetitions, \
        pc_query_ckDTree_k5_time[leaf_counter]/inps.nr_of_repetitions/60))
33     pc_ckDTree_distance, pc_ckDTree_id = pc_query_ckDTree(pc_xyz_ckDTree_tree, \
        pc_xyz, k=5)
34     pc_query_ckDTree_k5_stats[leaf_counter, :, 0] = np.mean(pc_ckDTree_distance, \
        axis=1)
35     pc_query_ckDTree_k5_stats[leaf_counter, :, 1] = np.std(pc_ckDTree_distance, \

```

```

        axis=1)
36  pc_query_ckDTree_k5_stats[leaf_counter, :, 2] = \
    np.percentile(pc_ckDTree_distance, [25], axis=1)
37  pc_query_ckDTree_k5_stats[leaf_counter, :, 3] = \
    np.percentile(pc_ckDTree_distance, [50], axis=1)
38  pc_query_ckDTree_k5_stats[leaf_counter, :, 4] = \
    np.percentile(pc_ckDTree_distance, [75], axis=1)
39  pc_ckDTree_distance = None
40  pc_ckDTree_id = None
41
42  print('\tQuerying ckDTree with k=10 for all points (%dx)... \
    '%(inps.nr_of_repetitions), end='', flush=True)
43  wrapped = wrapper(pc_query_ckDTree, pc_xyz_ckDTree_tree, pc_xyz, k=10)
44  pc_query_ckDTree_k10_time[leaf_counter] = timeit.timeit(wrapped, \
    number=inps.nr_of_repetitions)
45  print('time (average of %d runs): %0.3fs or %0.2fm'%(inps.nr_of_repetitions, \
    pc_query_ckDTree_k10_time[leaf_counter]/inps.nr_of_repetitions, \
    pc_query_ckDTree_k10_time[leaf_counter]/inps.nr_of_repetitions/60))
46  pc_ckDTree_distance, pc_ckDTree_id = pc_query_ckDTree(pc_xyz_ckDTree_tree, \
    pc_xyz, k=10)
47  pc_query_ckDTree_k10_stats[leaf_counter, :, 0] = np.mean(pc_ckDTree_distance, \
    axis=1)
48  pc_query_ckDTree_k10_stats[leaf_counter, :, 1] = np.std(pc_ckDTree_distance, \
    axis=1)
49  pc_query_ckDTree_k10_stats[leaf_counter, :, 2] = \
    np.percentile(pc_ckDTree_distance, [25], axis=1)
50  pc_query_ckDTree_k10_stats[leaf_counter, :, 3] = \
    np.percentile(pc_ckDTree_distance, [50], axis=1)
51  pc_query_ckDTree_k10_stats[leaf_counter, :, 4] = \
    np.percentile(pc_ckDTree_distance, [75], axis=1)
52  pc_ckDTree_distance = None
53  pc_ckDTree_id = None
54
55  print('\tQuerying ckDTree with k=50 for all points (%dx)... \
    '%(inps.nr_of_repetitions), end='', flush=True)
56  wrapped = wrapper(pc_query_ckDTree, pc_xyz_ckDTree_tree, pc_xyz, k=50)
57  pc_query_ckDTree_k50_time[leaf_counter] = timeit.timeit(wrapped, \
    number=inps.nr_of_repetitions)
58  print('time (average of %d runs): %0.3fs or %0.2fm'%(inps.nr_of_repetitions, \
    pc_query_ckDTree_k50_time[leaf_counter]/inps.nr_of_repetitions, \
    pc_query_ckDTree_k50_time[leaf_counter]/inps.nr_of_repetitions/60))
59
60  print('\tQuerying ckDTree with k=100 for all points (%dx)... \
    '%(inps.nr_of_repetitions), end='', flush=True)
61  wrapped = wrapper(pc_query_ckDTree, pc_xyz_ckDTree_tree, pc_xyz, k=100)
62  pc_query_ckDTree_k100_time[leaf_counter] = timeit.timeit(wrapped, \
    number=inps.nr_of_repetitions)
63  print('time (average of %d runs): %0.3fs or %0.2fm'%(inps.nr_of_repetitions, \
    pc_query_ckDTree_k100_time[leaf_counter]/inps.nr_of_repetitions, \
    pc_query_ckDTree_k100_time[leaf_counter]/inps.nr_of_repetitions/60))
64
65  print('\tQuerying ckDTree with k=500 for all points (%dx)... \
    '%(inps.nr_of_repetitions), end='', flush=True)

```

```

66 wrapped = wrapper(pc_query_cKDTree, pc_xyz_cKDTree_tree, pc_xyz, k=500)
67 pc_query_cKDTree_k500_time[leaf_counter] = timeit.timeit(wrapped, \
    number=inps.nr_of_repetitions)
68 print('time (average of %d runs): %0.3fs or %0.2fm'%(inps.nr_of_repetitions, \
    pc_query_cKDTree_k500_time[leaf_counter]/inps.nr_of_repetitions, \
    pc_query_cKDTree_k500_time[leaf_counter]/inps.nr_of_repetitions/60))
69
70 print('\tQuerying cKDTree with k=1000 for all points (%dx)... \
    '%(inps.nr_of_repetitions), end='', flush=True)
71 wrapped = wrapper(pc_query_cKDTree, pc_xyz_cKDTree_tree, pc_xyz, k=1000)
72 pc_query_cKDTree_k1000_time[leaf_counter] = timeit.timeit(wrapped, \
    number=inps.nr_of_repetitions)
73 print('time (average of %d runs): %0.3fs or %0.2fm'%(inps.nr_of_repetitions, \
    pc_query_cKDTree_k1000_time[leaf_counter]/inps.nr_of_repetitions, \
    pc_query_cKDTree_k1000_time[leaf_counter]/inps.nr_of_repetitions/60))
74 pc_xyz_cKDTree_tree = None
75 leaf_counter += 1
76
77 pc_cKDTree_time_df = pd.DataFrame({'index': range(len(leafrange)), 'leafsize': \
    leafrange,
78 'pc_query_cKDTree_k5_time': pc_query_cKDTree_k5_time.ravel(),
79 'pc_query_cKDTree_k10_time': pc_query_cKDTree_k10_time.ravel(),
80 'pc_query_cKDTree_k50_time': pc_query_cKDTree_k50_time.ravel(),
81 'pc_query_cKDTree_k100_time': pc_query_cKDTree_k100_time.ravel(),
82 'pc_query_cKDTree_k500_time': pc_query_cKDTree_k500_time.ravel(),
83 'pc_query_cKDTree_k1000_time': pc_query_cKDTree_k1000_time.ravel(),
84 'pc_generate_cKDTree_time': pc_generate_cKDTree_time.ravel()})
85 pc_cKDTree_time_df.to_hdf('pc_cKDTree_time_df_%s.hdf'%(inps.cpuname), \
    key='pc_cKDTree_time_df', complevel=9)
86
87 #For simplicity, save array to npy
88 np.save('pc_query_cKDTree_k5_stats_%s.npy'%(inps.cpuname), pc_query_cKDTree_k5_stats)
89 np.save('pc_query_cKDTree_k10_stats_%s.npy'%(inps.cpuname), pc_query_cKDTree_k10_stats)

```

Note that we also save the results to a pandas dataframe and the results of the queries for k=5 and k=10 to a .npy file.

This is repeated for *pyKDTree* (see Python code in github).

The generation of figures is outlined in `python/plot_py_and_cKDTree_variouscores.py` and